

Lecture 3 : Statistics for Human Geneticists



Volos Summer School

21 / 05 / 2018

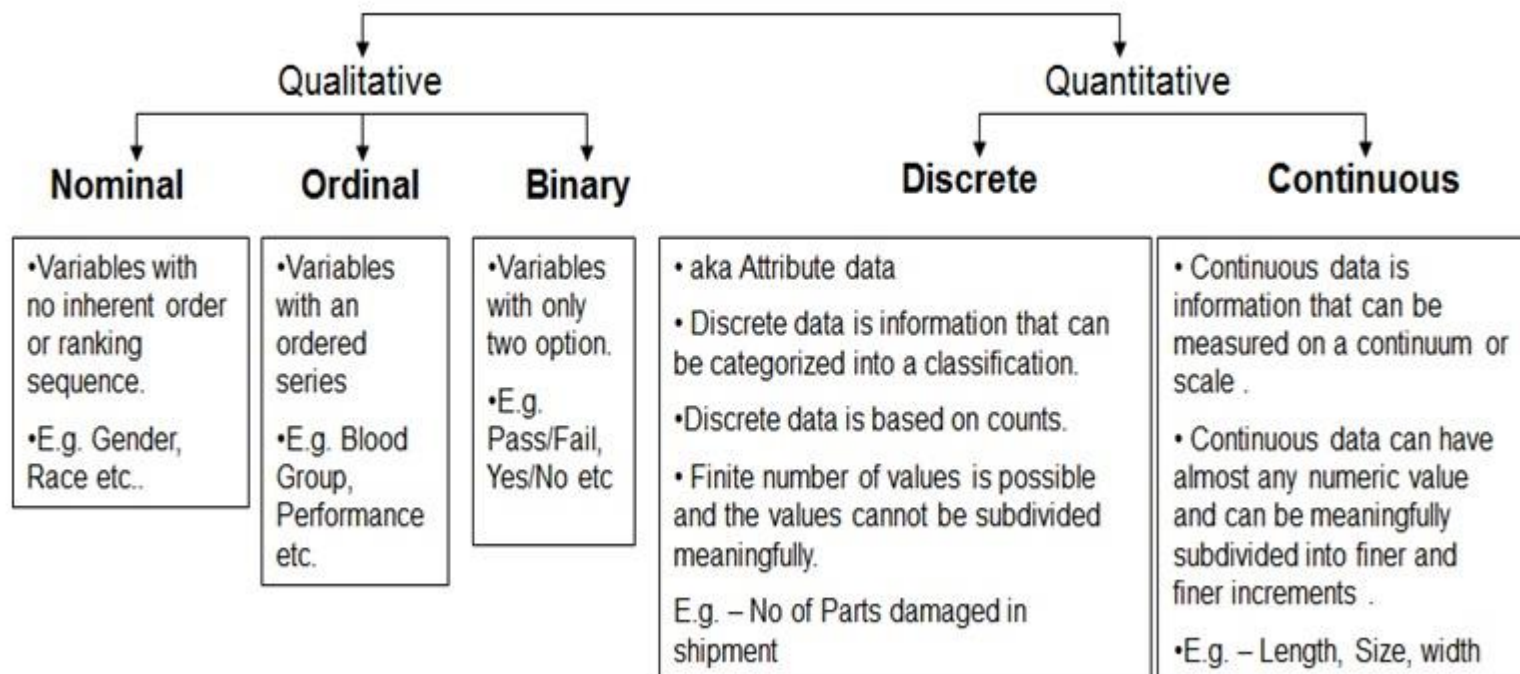
Arthur Gilly

What can we do with statistics?

- Estimation
- Hypothesis testing
- Modelling
- Predicting


Random Variables

- In statistics, we measure realizations of random variables
- Often, random variables follow a distribution
- They can be qualitative or quantitative, continuous or discrete




Distributions

- Two ways to represent them :



Cumulative distribution function (CDF)

- $y = p(X \leq x)$
- Always growing
- Ideal way to represent, but hard to read
- All distributions look the same




Probability density function (PDF)

- $y = p(X = x)$ for discrete
- Shows how values are distributed
- Nice visually, but mathematically hard to deal with


Distributions

- How to estimate them:



Empirical CDF

- Rarely used
- Except when you want to compute empirical quantiles



Barplot (discrete)


- For every value, count occurrences

Histogram (continuous)


- Cut the interval into bins, count observations within bin

Distributions

- Two broad types:
 - those followed by random variables (real world data)




$$X \sim \text{Poisson}(\lambda)$$




$$X \sim \mathcal{N}(\mu, \sigma^2)$$

Distributions

- Two broad types:
 - those followed by random variables (real world data)




$$X \sim \text{Poisson}(\lambda)$$




$$X \sim \mathcal{N}(\mu, \sigma^2)$$

Distributions

- Two broad types:
 - those followed by test statistics



$$X \sim T(\nu)$$



$$X \sim \chi^2(k)$$

λ , μ , σ , ν and k are ideal parameters. How to estimate them?

Statistics

- A statistic is a meaningful quantity derived from the data
- Often, estimators are realization of distribution parameters
- Examples? Mean, proportion
- For simple distributions/parameters, there is a formula
- For more complex ones, we have to use other techniques (Monte-Carlo, Permutations...)

$$(\hat{\mu} =) \bar{x} = \frac{1}{n} \sum_{k=0}^n x_k$$


$$\hat{p} = \frac{x}{n}$$

$$w = \frac{(\hat{\theta} - \theta_0)^2}{se(\hat{\theta})} \sim \mathcal{N}(0,1)$$

$$(\widehat{\sigma^2} =) s^2 = \frac{1}{N-1}$$

One particular statistic: Confidence intervals

- $x\%$ confidence interval ($x\%C.I.$) : $x\%$ of the time when this interval is calculated, it will contain the true value of the parameter



Hypothesis testing

- We want to measure whether the data gives sufficient evidence to reject a hypothesis
- Null/Alternative hypothesis ($\mathcal{H}_0/\mathcal{H}_A$)
- We prove that we can produce a statistic that follows a certain distribution if the null hypothesis is true = name of the test
- We calculate the statistic based on our data
- Because we know the distribution, we can calculate the CDF $p(X \leq x)$
- = how unlikely it is that our measurement comes from the null : p-value
- Example: proportion test, t-test, chi-squared test...

Summary statistic
(helps distinguish H_0 and H_A)



Test statistic
(standard distribution with no unknown parameters under H_0)




P-value
(probability of more 'extreme' test statistic)

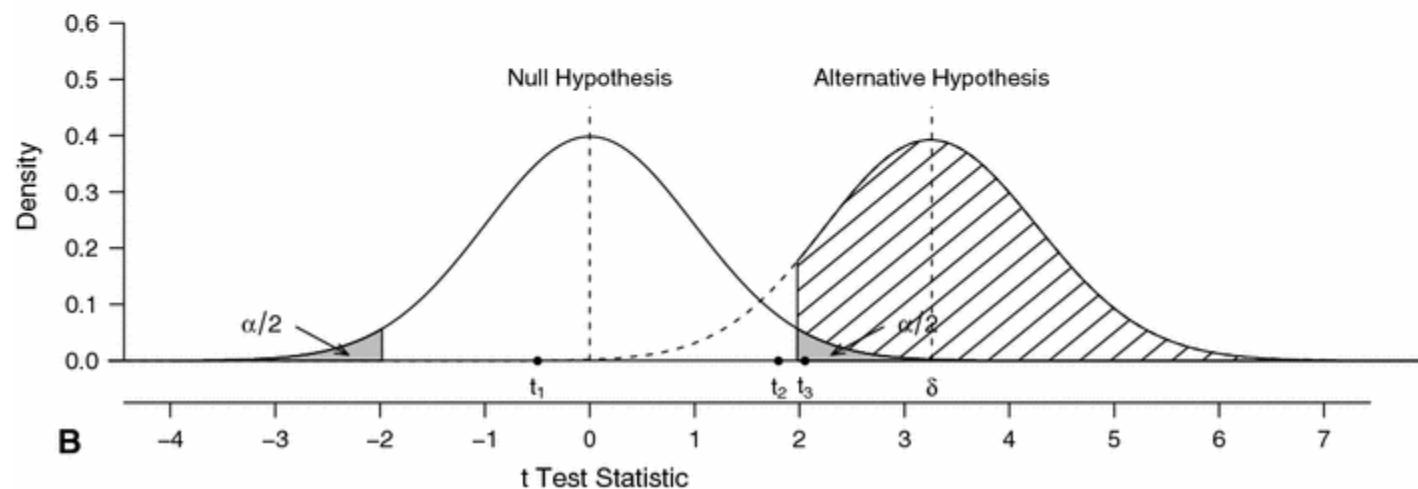

$$\chi^2 = \sum \frac{(n_{xy} - e_{xy})^2}{e_{xy}}$$




$\chi^2 \sim$ chi-squared $((r-1)(c-1) \text{ df})$



Hypothesis testing



Two tails or one



Exercise I : Proportion test

- In a population, we observe 41,009 potentially damaging variants among 14,281,180 variants
- What is the proportion? In a very large reference population, we observe a proportion of 1.52×10^{-3} . Is it significantly different? (prop.test, binom.test)

Multiple testing

Family-wise (FWER)

- Bonferroni correction
- Simple to implement, harder to interpret

$$p_{critical} = \frac{0.05}{m}$$

- “If all tests are under the null, probability that **one or more** of them is a false positive.”

False-discovery based (FDR)

- Benjamini-Hochberg procedure
- Harder to implement, easy to understand

$$p_{critical} = \operatorname{argmax}(p < \frac{i}{m} Q)$$

- i =rank, Q =FDR.
- “Proportion of significant tests that are false positives.”

Multiple testing

Family-wise (FWER)

- Bonferroni correction
- Simple to implement, harder to interpret

$$p_{critical} = \frac{0.05}{m}$$

- “If all tests are under the null, probability that **one or more** of them is a false positive.”

False-discovery based (FDR)

- Benjamini-Hochberg procedure
- Harder to implement, easy to understand

$$p_{critical} = \operatorname{argmax}(p < \frac{i}{m} Q)$$

- i =rank, Q =FDR.
- “Proportion of significant tests that are false positives.”

When to use which depends on

- 1) best practices
- 2) relative price of false negative/positive

Checking results

- Statistical significance:
 - One test: $p < 0.05$
 - Genome-wide: one test per variant and per phenotype
 - But all variants are not independent, in reality, we account for LD
 - 5×10^{-8} for GWAS, 10^{-9} for sequencing-based

Dietary variable	<i>P</i> value
Total calories	<0.001
Olive oil	0.008
Whole milk	0.039
White meat	0.041
Proteins	0.042
Nuts	0.060
Cereals and pasta	0.074
White fish	0.205
Butter	0.212
Vegetables	0.216
Skimmed milk	0.222
Red meat	0.251
Fruit	0.269
Eggs	0.275
Blue fish	0.34
Legumes	0.341
Carbohydrates	0.384
Potatoes	0.569
Bread	0.594
Fats	0.696
Sweets	0.762
Dairy products	0.94
Semi-skimmed milk	0.942
Total meat	0.975
Processed meat	0.986

Dietary variable	<i>P</i> value	Rank	(i/m)Q
Total calories	<0.001	1	0.010
Olive oil	0.008	2	0.020
Whole milk	0.039	3	0.030
White meat	0.041	4	0.040
Proteins	0.042	5	0.050
Nuts	0.060	6	0.060
Cereals and pasta	0.074	7	0.070
White fish	0.205	8	0.080
Butter	0.212	9	0.090
Vegetables	0.216	10	0.100
Skimmed milk	0.222	11	0.110
Red meat	0.251	12	0.120
Fruit	0.269	13	0.130
Eggs	0.275	14	0.140
Blue fish	0.34	15	0.150
Legumes	0.341	16	0.160
Carbohydrates	0.384	17	0.170
Potatoes	0.569	18	0.180
Bread	0.594	19	0.190
Fats	0.696	20	0.200
Sweets	0.762	21	0.210
Dairy products	0.94	22	0.220
Semi-skimmed milk	0.942	23	0.230
Total meat	0.975	24	0.240
Processed meat	0.986	25	0.250

Modelling and predicting

- If we estimate the effect of one variable on another variable, we do modelling
- When we apply this effect to new observations of the variable, we do prediction
- Process is called machine learning, predictive modelling or predictive analysis
- In human genetics, main task is to model effect of genotypes on phenotypes

$$\textit{phenotype} \sim \beta \times \textit{genotype} + \epsilon$$

$$\begin{bmatrix} \textit{pheno}_0 \\ \vdots \\ \textit{pheno}_n \end{bmatrix}$$

$$\begin{bmatrix} A/T \\ \vdots \\ T/T \end{bmatrix}$$

Modelling and predicting

- If we estimate the effect of one variable on another variable, we do modelling
- When we apply this effect to new observations of the variable, we do prediction
- Process is called machine learning, predictive modelling or predictive analysis
- In human genetics, main task is to model effect of genotypes on phenotypes

$$\textit{phenotype} \sim \beta \times \textit{genotype} + \epsilon$$

$$\begin{bmatrix} \textit{pheno}_0 \\ \vdots \\ \textit{pheno}_n \end{bmatrix}$$

= {0,1} (case-control)
 $\in \mathbb{R}$ (quantitative) $\sim \mathcal{N}(0,1)$

$$\begin{bmatrix} A/T \\ \vdots \\ T/T \end{bmatrix}$$

= {0,1,2} (genotype, directly typed)
 $\in [0,2]$ (dosage, imputed)

$$\begin{bmatrix} 1 \\ \vdots \\ 2 \end{bmatrix}$$

$$\begin{bmatrix} 0.965 \\ \vdots \\ 1.816 \end{bmatrix}$$

Modelling and predicting

- If we estimate the effect of one variable on another variable, we do modelling
- When we apply this effect to new observations of the variable, we do prediction
- Process is called machine learning, predictive modelling or predictive analysis
- In human genetics, main task is to model effect of genotypes on phenotypes

Usually, we do not predict (except PRS)

$$\textit{phenotype} \sim \beta \times \textit{genotype} + \epsilon$$

$$\begin{bmatrix} \textit{pheno}_0 \\ \vdots \\ \textit{pheno}_n \end{bmatrix} \quad \begin{bmatrix} A/T \\ \vdots \\ T/T \end{bmatrix} = \begin{bmatrix} 1 \\ \vdots \\ 2 \end{bmatrix}$$




$= \{0,1\}$ (case-control)
 $\in \mathbb{R}$ (quantitative) $\sim \mathcal{N}(0,1)$

$= \{0,1,2\}$ (genotype, directly typed)
 $\in [0,2]$ (dosage, imputed)

$$\begin{bmatrix} 0.965 \\ \vdots \\ 1.816 \end{bmatrix}$$

Case/control

- Estimated effect: odds ratio (OR)
“how much more likely are you to be a case if you carry the risk allele?”
 per genotype, calculate the odds $O = \frac{p}{1-p}$





Case/control


- Estimated effect: odds ratio (OR)

“how much more likely are you to be a case if you carry the risk allele?”

per genotype g and for a disease Y , calculate the odds $O = \frac{p_{Y=1|g}}{1-p_{Y=1|g}}$



	cases	controls
T	380	490
C	620	510



$$O_T = \frac{380/n_T}{490/n_T} \quad O_C = \frac{620/n_C}{510/n_C}$$

$$OR_{C/T} = \frac{620 \times 490}{510 \times 380} = 1.56$$

Case/control

Dominant

Marker allele	Affected	Unaffected
DD+Dd	$n_{2A} + n_{1A}$	$n_{2U} + n_{1U}$
dd	n_{0A}	n_{0U}

$$OR = \frac{n_{affected\ carriers} \times n_{healthy\ non-carriers}}{n_{healthy\ carriers} \times n_{affected\ non-carriers}}$$

Recessive

Marker allele	Affected	Unaffected
DD	n_{2A}	n_{2U}
Dd+dd	$n_{1A} + n_{0A}$	$n_{1U} + n_{0U}$

Additive

Marker genotype	Affected	Unaffected
DD	n_{2A}	n_{2U}
Dd	n_{1A}	n_{1U}
dd	n_{0A}	n_{0U}

$$OR = \frac{(2 \times n_{2A} + n_{1A}) \times (2 \times n_{0U} + n_{1U})}{(2 \times n_{0A} + n_{1A}) \times (2 \times n_{2U} + n_{1U})}$$


Allelic odds-ratio

Case/control


- Output: OR and 95% confidence interval of the OR
- Test: is it significantly different from 1?
- Tests: Fisher's exact test or Chi-squared
- In case of dosages or covariates: logistic regression

Continuous trait

- For directly typed (0,1,2):ANOVA





additive



recessive

Continuous trait

- For dosages (imputed quantity of minor allele $d \in [0,1]$) : linear regression
- In general: generalized linear model



Continuous trait

A linear regression model is defined as

$$y = x\beta_1 + \beta_0 + \varepsilon$$

Data:


- y : a continuous trait
- x : SNP genotype at a given locus

Parameters:

- β_1 : regression coefficient, represents the strength of association between x and y
- β_0 : intercept term (is 0 or ignored)
- ε : noise or the part of y that is not explained by x (e.g., environmental effect)


Assumptions:

- The individuals in the study are not related
- The phenotype y has a normal distribution





Checking results

- QQ-plot
 - Distribution of p-values is uniform $[0,1]$ under the null
 - If we have much signal, more around 0
 - Compare quantiles with expected ones : QQ-plot
 - In R: qqunif



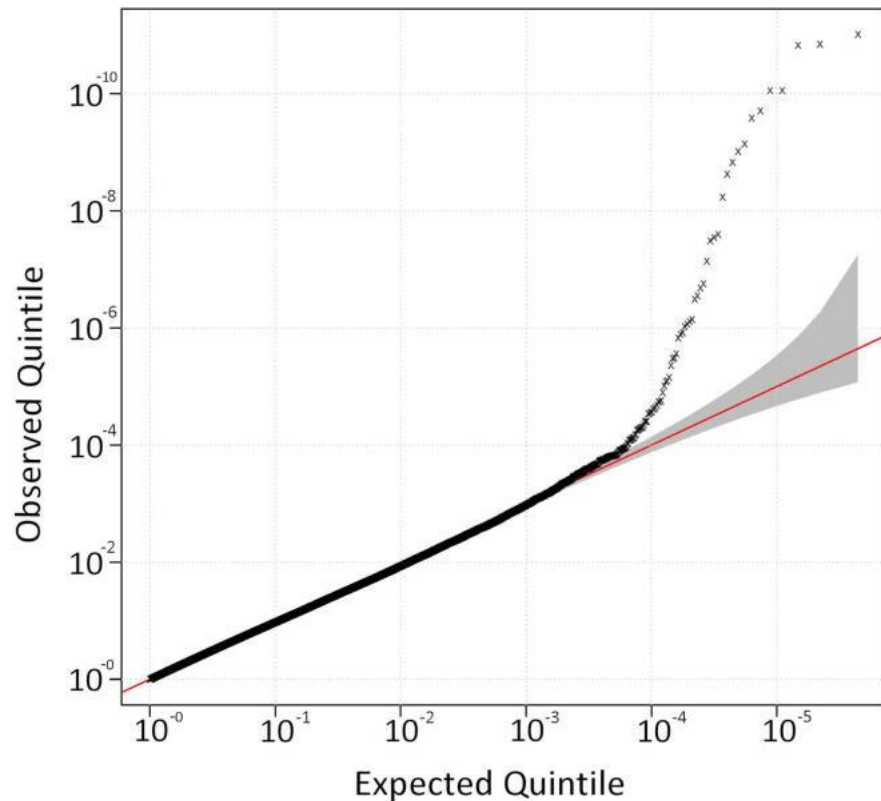
Checking results

- QQ-plot
 - Distribution of p-values is uniform $[0,1]$ under the null
 - If we have much signal, more around 0
 - Compare quantiles with expected ones : QQ-plot
 - In R: qqunif



Checking results


- QQ-plot
 - Distribution of p-values is uniform $[0,1]$ under the null
 - If we have much signal, more around 0
 - Compare quantiles with expected ones : QQ-plot
 - In R: qqunif



- Inflation: too much signal
- Measured visually, but also lambda

Checking results

- QQ-plot
 - Distribution of p-values is uniform $[0,1]$ under the null
 - If we have much signal, more around 0
 - Compare quantiles with expected ones : QQ-plot
 - In R: qqunif



Appearances can be deceiving:


- A QQ-plot can look inflated when it isn't (just a lot of signal)
- And conversely
- We calculate the genomic inflation factor

$$\lambda = \frac{\text{median}(Q_{\chi^2}(p))}{0.45}$$

(median of χ^2 test statistics divided by median of χ_1^2)

Checking results

- QQ-plot
 - Distribution of p-values is uniform $[0,1]$ under the null
 - If we have much signal, more around 0
 - Compare quantiles with expected ones : QQ-plot
 - In R: qqunif



Appearances can be deceiving:

- A QQ-plot can look inflated when it isn't (just a lot of signal)
- And conversely
- We calculate the genomic inflation factor

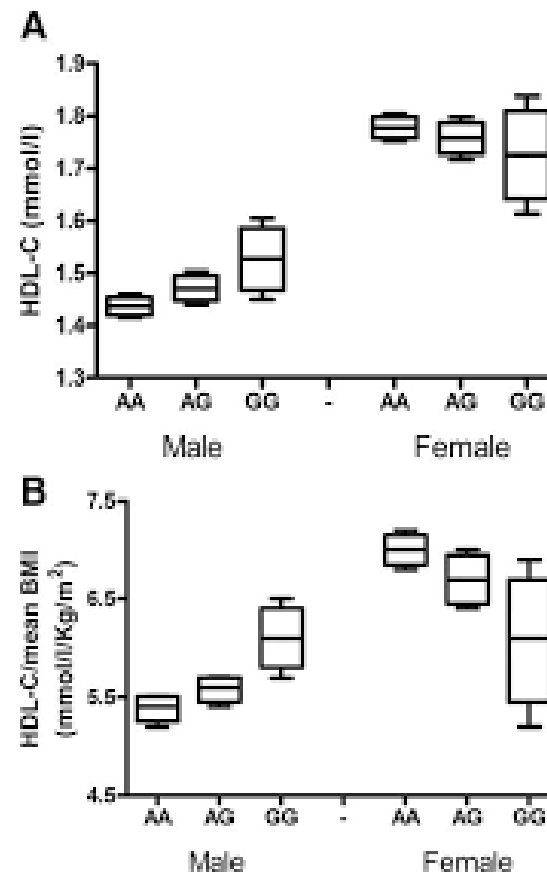
$$\lambda = \frac{\text{median}(Q_{\chi^2}(p))}{0.45}$$

(median of χ^2 test statistics divided by median of χ^2_1)

- Ideally, want to correct in the model
- Can also adjust: GC correction (divide by lambda)

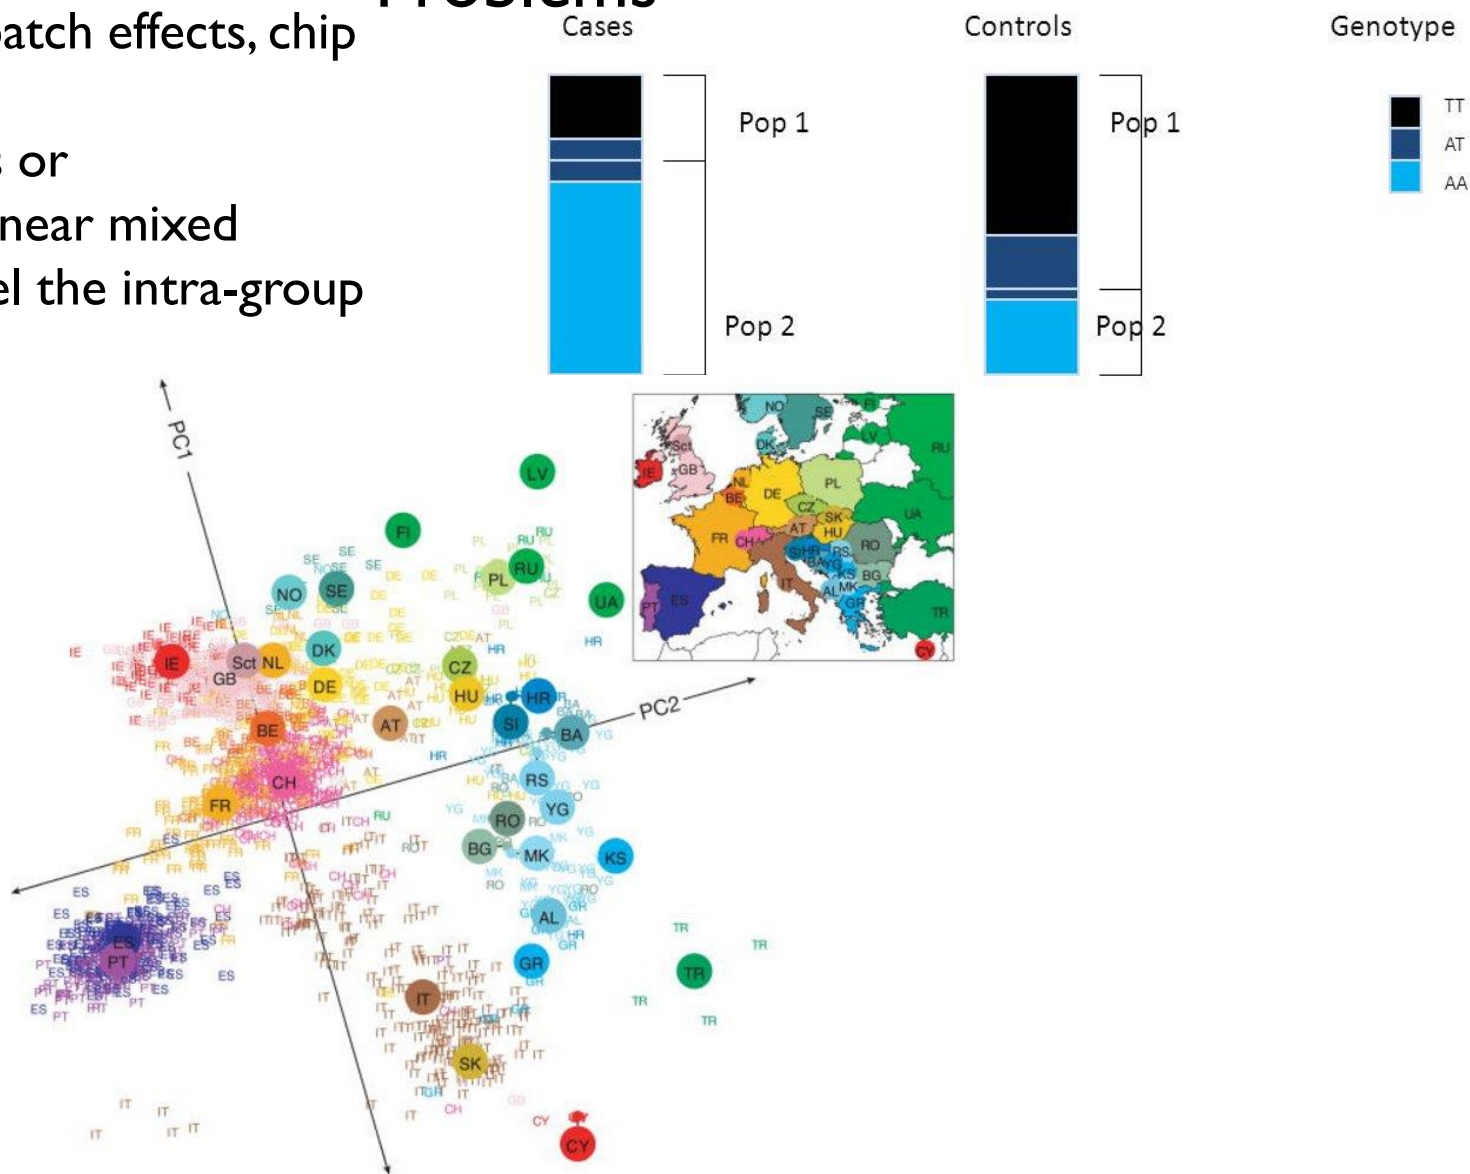
Problems

- Covariates: sex, batch effects, chip effects



Problems



- Covariates: sex, batch effects, chip effects
- Structure: villages or subpopulations: linear mixed models can model the intra-group effect



Problems


- Covariates: sex, batch effects, chip effects
- Structure: villages or subpopulations: linear mixed models can model the intra-group effect

$$\text{phenotype} \sim \beta \times \text{genotype} + \beta_1 \times \text{covariates} + \beta_2 \times \text{structure} + \epsilon$$



References

*These references are for probability and statistical theory only
They are not specific to statistical genetics*



- Murphy K, **Machine learning: A probabilistic perspective.** (MIT Press)



- Johnson and Wichern, **Applied multivariate statistical analysis.** (Pearson)